

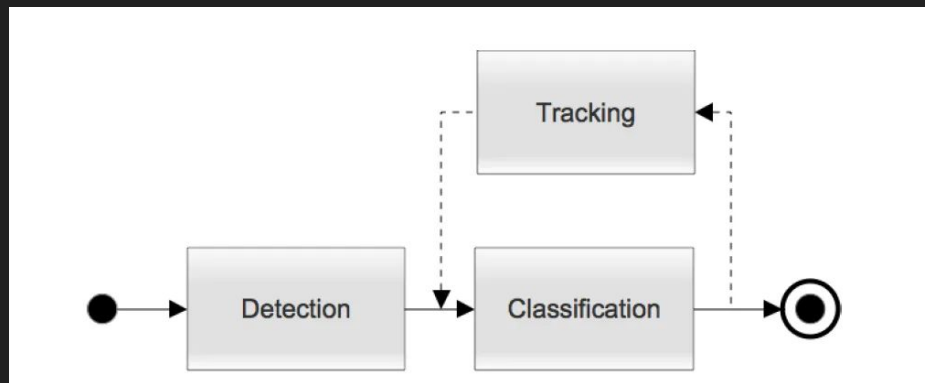
# CDR – 2A1

Traffic Light Detection and Tracking

Morgan Roberts, Aaryan Shenoy, Clayton Gowan,  
Xiaohu (Max) Huang, Robert Madriaga

# Project Background

# Background



## The Problem

- Fast and reliable detection of traffic light, and light states
- Can be broken down into 3 parts
  - Detection
  - Classification
  - Tracking (loops back to classification)

# Background

## Traffic Light Detection Requirements

- Detection
  - Identify image
  - Deep Learning is a natural solution
    - Can be computationally expensive
- Classification
  - Lights must be classified by:
    - Type (arrow direction, solid, blinking, etc)
    - Color (red, yellow, green)
- Tracking
  - Faster than detection/scanning of the whole image
  - Conserve data between frames
    - Compensates for lack of data from detection



# Needs Statement


# Needs Statement

- Need – to detect, track and classify traffic lights
  - Prioritize speed and accuracy
  - Track lights across multiple frames
- Why – Self-driving vehicles must react to traffic lights.
  - Many systems use image data for this purpose.
  - These algorithms are usually too slow.
    - Reaction times are vital for safety

# Goals and Objectives

# Goals and Objectives

- Goal: Develop a speed-focused detection model for traffic light types and colors, analyze the states of the lights, track detection results across multiple frames, and integrate with ROS.
- Motivation:
  - Fast detection and tracking speeds are crucial in autonomous driving applications
  - Modularization of detection and tracking components





# Design Alternatives



# Multi-Object Trackers

- DeepSORT
  - Popular tracking algorithm that utilizes frame trajectory predictions and deep learning
  - Uses YOLOv8 bounding boxes to predict location in successive frames
- Norfair
  - Open-source Python library that supports object tracking in moving cameras
  - Compatible with YOLOv8 model through use of its bounding boxes
  - Use of Norfair is promising compared to DeepSORT
    - Good Performance
    - Ease of implementation
    - DeepSORT issues

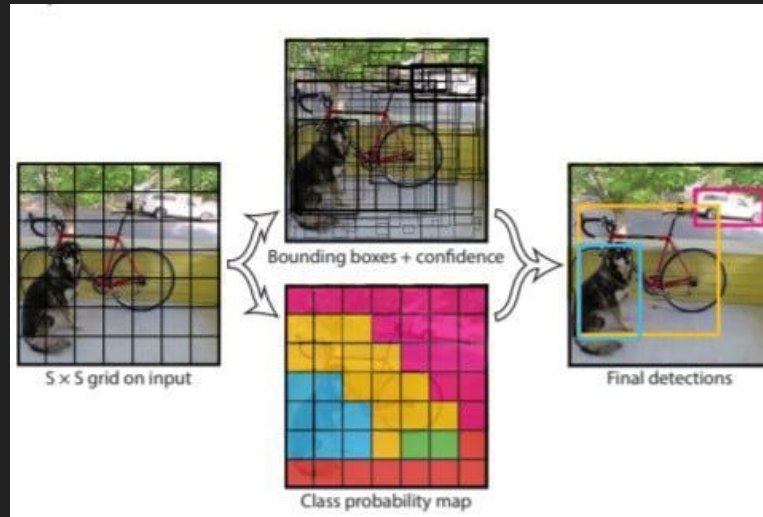
# Proposed Design

# Proposed Design

- Input: ROS Bag Datasets
- Output: Camera frames with detected traffic lights labeled
- Tools to Use:
  - Machine learning framework
    - YOLOv8 detection model
    - Norfair library
  - ROS Bag Screenshot Dataset
  - ROS node system

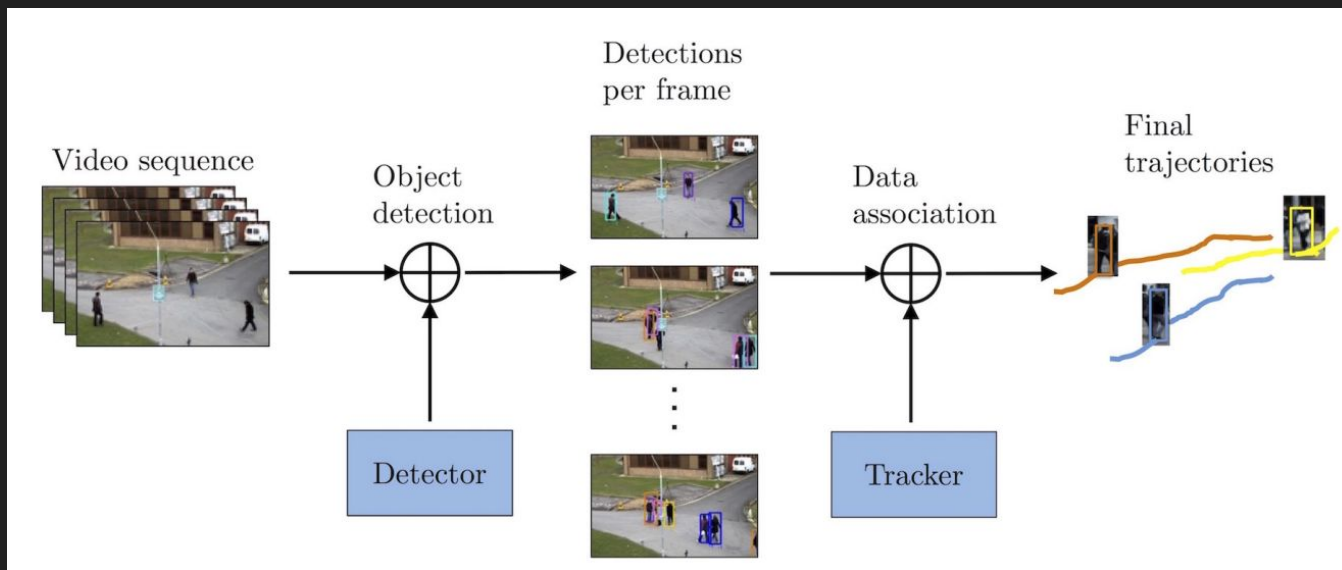
# Proposed Design

- YOLOv8 detection model
  - Runs on each camera frame to produce detections



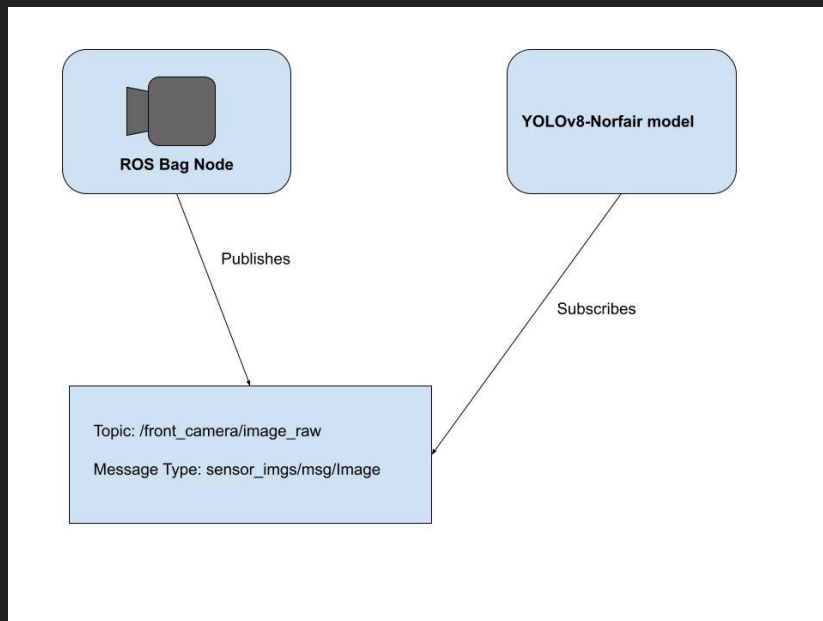
# Proposed Design

- Norfair Python library
  - Library functions utilize YOLOv8 bounding boxes to track traffic lights across frames



# Proposed Design

- ROS and machine learning integration via publisher-subscriber model





Vision for Resulting Product


# Current Progress

- ROS node network is set up
  - Contains processes that play camera footage and run YOLO and object tracking models
- Python program created to process each input frame through YOLO and object trackers
- YOLO/Norfair model bounding boxes detects lights better than YOLO/DeepSORT due to default DeepSORT feature extractor trained on humans

# Current Progress, Cont.



Detection



Tracking

Note: Video is played at 3x actual speed


# End Result Vision

- Take camera feed as input
- Have the YOLO model detect and classify traffic lights accurately and quickly
- Use detections to continually update the tracker, which will assign IDs
- Filter the tracked objects to only include the relevant traffic lights
- Determine flashing light states
- Output the currently tracked traffic lights (amount, type, etc.) in the desired output

# Approach for Design Validation

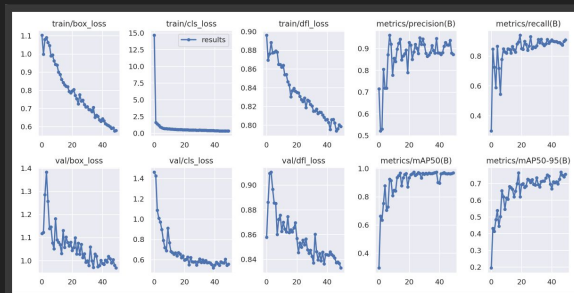
# Approach for Design Validation

- Ensure we are running at an adequate framerate with good detection speed while also maintaining a high level of confidence and accuracy in detection
  - Measurements: FPS, detection model confidence levels
- Find an optimal balance between model size, speed, and accuracy
  - Compare pre-trained models/different datasets with custom models
- Verify that multiple objects can be tracked accurately across frames
- Make sure pipeline functions correctly with desired formats/outputs



# Approach for Design Validation, Cont.

- Utilize YOLO training output to compare model performance on testing and validation data, observe inference speed and performance on rosbag files
- mAP, Recall, etc.



```
Epoch GPU_mem box_loss cls_loss dfl_loss Instances Size
50/50 13.8G 0.5782 0.3309 0.7987 46 1280: 100% 46/46 [01:10<00:00, 1.53s/it]
Class Images Instances Box(P R mAP50 mAP50-95): 100% 3/3 [00:09<00:00, 3.21s/it]
all 69 176 0.873 0.908 0.969 0.756
```

50 epochs completed in 2.149 hours.

Optimizer stripped from runs/detect/train/weights/last.pt, 22.7MB

Optimizer stripped from runs/detect/train/weights/best.pt, 22.7MB

Validating runs/detect/train/weights/best.pt...


Ultralytics YOLOv8.0.53 Python-3.9.16 torch-1.13.1+cu116 CUDA:0 (Tesla T4, 15102MiB)

Model summary (fused): 168 layers, 11128293 parameters, 0 gradients, 28.5 GFLOPs

Class	Images	Instances	Box(P)	R	mAP50	mAP50-95)
all	69	176	0.915	0.884	0.965	0.755
go	69	75	1	0.775	0.958	0.661
goLeftStop	69	4	0.914	1	0.995	0.895
stop	69	71	0.944	0.718	0.93	0.593
stopLeft	69	1	0.619	1	0.995	0.895
warning	69	20	0.996	0.8	0.885	0.73
warningLeft	69	2	0.933	1	0.995	0.798
warningLeftStop	69	3	1	0.895	0.995	0.708

Speed: 6.2ms preprocess, 28.2ms inference, 0.0ms loss, 2.0ms postprocess per image

Results saved to runs/detect/train



# Economic Analysis/Budget



# Economic Analysis/Budget Update

Each team gets a \$500 Budget, and out of our budget we spend:

Human Labor - \$0

Software - Open Source - \$0

Software - Closed Source / Cloud Services - \$variable


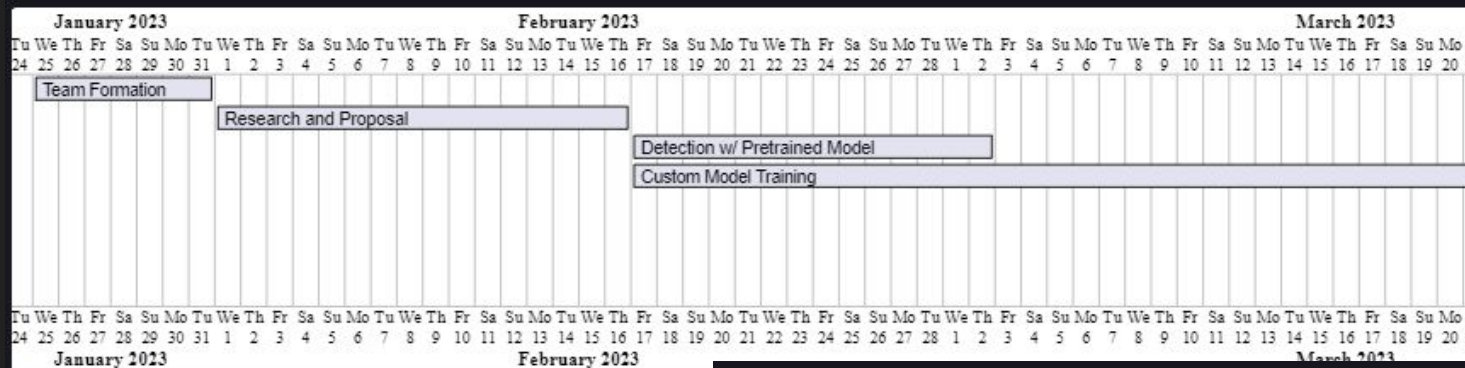
- Money could be spent on an upgraded Colab subscription
- This has not been necessary so far



# Schedule of Tasks

# Gantt Chart

- Estimation of development schedule (subject to change)
- Model Training took longer than expected



# Project Management and Teamwork


# Teamwork and Roles

- Morgan Roberts: Team Leader
  - Aaryan Shenoy: Systems Design
  - Clayton Gowan: Software Design and Model Training
  - Xiaohu (Max) Huang: Software Design and Testing
  - Robert Madriaga: Technical Writing and Dataset Management
- 
- There have been weekly in-person team meetings on Tuesdays/Thursdays and Discord meetings as needed
  - Members may assist in other roles/responsibilities as needed

# Societal, Safety, and Environmental Analysis

# Societal, Safety, and Environmental Analysis

- Autonomous vehicles have the potential to improve convenience and quality of life.
  - Optimal travel times
  - Increased accessibility to transportation
- Privacy Concerns
  - Requires monitoring equipment
- Safety Concerns
  - Even minute errors in software or hardware can be catastrophic
  - Lead to loss of life and property damage
- Safety Benefits
  - With ideal monitoring, computation, and decision making, a self driving car can react to and avoid dangerous situations faster than humans and with more consistency
- Environmental Benefits
  - Effective navigation leads to a reduction in emissions



# Manufacturability

- **Software Compatibility**

- Software is compatible with Linux operating systems that is able to run ROS2 Foxy and Python 3, or through a VM box or docker environment
- Lack of support or updates to any piece of software used may cause errors as features may be deprecated or incompatibility with newer versions

- **Hardware Compatibility**

- Designed to work with high resolution camera at 60 fps
- Requires moderately powerful modern hardware to run on
- Higher frame rate cameras may cause issues with blinking light detection code



Questions